MIDS W205

Lab #	6	Lab Title	Apache Spark SQL - a few features
Related Module(s)	6	Goal	Get you started on spark and pyspark
Last Updated	9/24/15	Expected duration	20 to 30 minutes

Todos:
- Improve introduction
- Explain what they will learn.
- For each step explain what and why and how to verify
- Include references to useful resources.

Install spark
Run simple spark command using the pyspark shell
Use ipython
Loading data
Run a simple command
      - count words
Use spark-sql

This Lab

In this lab you will get an hands on introduction to Sparks and specifically to Spark SQL. After this lab you should be able to launch Spark, create tables, load data and run simple queries using Spark SQL. We will go over Apache Spark features and common commands as below:

1. Accessing Spark SQL - CLI
2. Accessing through Java database connectivity (JDBC)
3. Accessing through Java database connectivity - Beenline
4. Let's create a table and load data using CSV file
5. Accessing Spark-SQL in Python Code
6. Caching tables and Un-caching tables

Intro

Apache Spark
Apache Spark is an open-source distributed computing. This was developed in the AMPLab at UC Berkley. Spark uses in-memory processing and can sometime be up to 100 times faster than regular Hadoop Map Reduce on certain types of jobs.

Apache Spark has a basic computing substrate and several frameworks built on top of it. There is a framework for querying structured data (Spark SQL), a streaming analytics framework, a machine learning framework etc.

Spark SQL is a Spark module in which uses structured data processing. It provides a programming abstraction called DataFrames. Spark SQL acts as a distributed SQL query engine.

Here are some useful resources:

Resource	What
http://spark.apache.org/docs/latest/programming-guide.html	This guide shows each of these features in each of Spark's supported languages. It is easiest to follow along with if you launch Spark's interactive shell. This includes reference to the basic commands you can perform on RDD's. Such as filter records, count records, join records.
https://spark.apache.org/docs/1.1.0/sql-programming-guide.html	Guide for using Spark SQL.
https://spark.apache.org/docs/0.9.0/python-programming-guide.html	Python spark programming guide.
http://spark.apache.org/docs/latest/api/scala/index.html#org.apache.spark.SparkContext	Programming guide for the Spark Context object. Here you can find actions available on the Spark Contexts.

In Spark, a DataFrame is a distributed collection of data organized into named columns. It is conceptually equivalent to a table in a relational database or a data frame in R/Python, but with richer optimizations. DataFrames can be constructed from a wide array of sources such as: structured data files, tables in Hive, external databases, or existing RDDs.

Here is an example which shows how to construct DataFrames in Python programming language :

```
//Constructs a DataFrame from the users table in Hive.
users = context.table("users")
//Constructs a DataFrame from from JSON files in S3
logs = context.load("s3n://path/to/data.json", "json")
```

Here are a few supported Data Formats and Sources :

1. JSON files, Parquet files, Hive tables
2. Local file systems
3. Distributed file systems (HDFS)
4. Cloud storage (S3)
5. External relational database systems via JDBC
6. Extensions include Avro, CSV, ElasticSearch, and Cassandra

Step-1.Accessing Apache Spark SQL (CLI)

Using the command line interface (CLI)

From Spark Installation folder, let's go to spark-sql prompt. Please check where is your spark application is installed.

```
cd /usr/lib/spark

./bin/spark-sql
```

Once you are in spark prompt (spark-sql) , to list the tables, you can run as follows:

```
spark-sql> show tables;
```

If you want to connect via spark-shell, here is a very useful link:
https://docs.sigmoidanalytics.com/index.php/Interactive_Analysis_with_the_Spark_Shell

Step-1. Check installation

Echo your SPARK_HOME environment variable to see where you Spark is installed.

```
%echo $SPARK_HOME
```

Got to that directory and look what is in it.

```
my-mac:lab_6 my$ cd $SPARK_HOME
my-mac:spark my$ ls
CHANGES.txt NOTICE       README.md   bin   data  ec2   lib   python
LICENSE     R            RELEASE     conf  derby.log  examples   old
sbin
```

If you look in bin you will see spark-shell, pyspark and other tools for running and managing spark.

You should see the following result

You should see something like:
```
jkoister-mac:lab_6 jkoister$ echo $SPARK_HOME
/Applications/devtools/spark
```

Step-1. Start pyspark

First we will start a spark shell so that we can access spark and interactively process sprak commands. We will be using pyspark, which is a python based shell for spark.

Start the shell
```
./pyspark
```

In the pyspark shell you can use python instructions. Create a variable with some value using this command.

```
>>> x = [1,2,3,4,5,6,7,8,9];
>>> print x
[1, 2, 3, 4, 5, 6, 7, 8, 9]
>>> len(x)
9
```

There is something call a spark context. The spark context is the "object" you use to refer to the spark cluster. type sc to verify you have one:

```
>>> sc
<pyspark.context.SparkContext object at 0x1063b3410>
```

So far you only used Python statements. You can use the Spark context to create a Spark RDD from this data using the following command:

```
>>> distData = sc.parallelize(x);
```

```
>>> print distData;
ParallelCollectionRDD[0] at parallelize at PythonRDD.scala:391
```

The Spark Context parallelize action is used to take local programming collections and create RDDs from them. In this case we created a RDD from a Python array.

distData is an RDD representation. Try doing len(distData), what happens and why? To count elements in an RDD you need to use RDD actions. You can find a list of actions in the programming guide. To count the elements you use the count() action. Try the following:

```
>>> nx=distData.count()
>>> print nx
9
```

The default level of logging can be distracting. To reduce the logging information go to the $SPARK_HOME/conf directory. Create a log4j.properties file. If you do not already have one you can do that by copying the log4j.properties.template file. Change the logging level to warnings only by change INFO to WARN in the property log4j.rootCategory=WARN, console

```
# Set everything to be logged to the console
log4j.rootCategory=WARN, console
log4j.appender.console=org.apache.log4j.ConsoleAppender
log4j.appender.console.target=System.err
log4j.appender.console.layout=org.apache.log4j.PatternLayout
```

If you rerun the commands above you should see much less output in the Spark shell.

You should see the following result

If you

Step-2. Load a file and count the rows.

Spark is commonly used to process large sets of data, and naturally we often read these data files from disc. The action textFile is an operation on Spark Contexts that creates a RDD from a file that resides in HDFS or the local files system. Make sure you cloned the data file Crimes_-_2001_to_present.csv in the lab directory. Now create an RDD from that file using this action:

```
crimedata =sc.textFile("Crimes_-_2001_to_present.csv")
```

Print the number of lines this RDD using the following command:
```
>>> print crimedata.count()
5862796
```

```
>>>
```

As you can see there are almost 6 million records, so it took a few seconds to count them.

You can get the first element of the RDD with the operation first:

```
>>> crimedata.first()
```

You can get the n first elements with the operation take(n)
```
>>> crimedata.take(10)
```

As you see the data include the header information. Remove the header from the RDD is not straight forward. The reason is that RDD is immutable by design. One way to remove it is to create a new RDD and to filter the first row out. There are many ways to do that, here is an example:

```
>>> noHeaderCrimedata = crimedata.zipWithIndex().filter(lambda
(row,index): index > 0).keys()
```

It is a quite a processing heavy way of doing it, but as we mentioned RDD's are immutable you can not just go in and remove a record. It is an inherent assumption to Spark which allows it to certain things more efficiently.

Print the first line to check that the header if gone. And count the lines to make sure the number seems correct.

```
>>> noHeaderCrimedata.first()
>>> noHeaderCrimedata.count()
```

Here is a more Python like way of doing the same, that may be easier to understand. We first define a Python function in our Python spark shell using the following line.

```
>>> def remove_header(itr_index, itr): return iter(list(itr)[1:]) if
itr_index == 0 else itr
```

We then execute a mapPartitionWithIndex operations passing that function as an argument.

```
>>> noHeaderCrimeData2 = crimedata.mapPartitionsWithIndex(remove_header)
```

Print the first line and the count to make sure it is correct.

You should see the following result

Step-3. Filter records, structures

One obvious operation for Spark is to filter the data. Lets filter out all crimes that seems to be related to "NARCOTICS".

We can do this using the filter operations and a lambda function that checks if the word "NARCOTICS" appears in the each row. We will only return rows that included that word.

```
narcoticsCrimes = noHeaderCrimedata.filter(lambda x: "NARCOTICS" in x)
>>> narcoticsCrimes.count()
663712
```

It appears that the are 663712 crimes related to narcotics. Use take(n) to check that the data seems ok. For example:

```
>>> narcoticsCrimes.take(20)
```

The RDD we have are just long strings, the fact that fields are comma separated does not mean anything to the Spark RDD. If we want to create a structure with which we want to do some more advanced things with we need to parse the rows and create the appropriate structure. The operations below splits each record up as an array using the Python operation split. It then create a new RDD were each row is an array of strings as opposed to one long string.

```
>>> narcoticsCrimeRecords = narcoticsCrimes.map(lambda r :
r.split(","))
```

You can see the first array record using:
```
>>> narcoticsCrimeRecords.first()
```

You can check that you still have the same number of rows using:
```
>>> narcoticsCrimeRecords.count()
```

You should see the following result

Step-3. Key-values

An important structure in spark is called Key Value pairs. In Python those are represented as Python tuples. A tuples is an immutable sequence of elements of various types.

You can create a new RDD consisting of tuples using the following operation:
```
>>> narcoticsCrimeTuples = narcoticsCrimes.map(lambda x:
(x.split(",")[0], x))
```

You can check that the number of tuples is the same as the number of records in the data.

```
>>> narcoticsCrimeTuples.count()
```

It takes the first first element and makes it a key and the rest if the row becomes the value part of the tuple. You can examine the tuple using the following operations:

```
>>> narcoticsCrimeTuples.first()
```

And you can check it out using these RDD and Python functions.
Get the first tuple:

```
>>> firstTuple=narcoticsCrimeTuples.first()
```

How many elements do you have in the tuple?

```
>>> len(firstTuple)
```

What is the key of the first tuple?

```
>>> firstTuple[0]
```

What is the value of the first tuple?

```
>>> firstTuple[1]
```

Step-4. Start Spark-SQL

You should see the following result

Step-5. Let's create a table and load data using CSV file:

Creating a table
You can create a Spark SQL table using the following format.

```
create table Web_Session_Log
(DATETIME varchar(500),
USERID varchar(500),
SESSIONID varchar(500),
PRODUCTID varchar(500),
REFERERURL varchar(500))
row format delimited fields terminated by '\t'
stored as textfile;
```

Now let's load data to this table from HDFS and Local:

HDFS :
```
LOAD DATA INFILE '/mnt/weblog.csv' INTO TABLE Web_Session_Log;
```

Local file system :
```
LOAD DATA LOCAL INFILE '/mnt/weblog.csv' INTO TABLE Web_Session_Log;
```

You, please create a datafile of your own choice. Also, create a table using Spark-sql and load the data.

You should see the following result

Step-6. Accessing Spark-SQL in Python Code:

Programming Spark SQL with Python

Here is a sample code:
```
from pyspark.sql import SQLContext
sqlContext = SQLContext(sc)
Running SQL Queries Programmatically :
from pyspark.sql import SQLContext
sqlContext = SQLContext(sc)
df = sqlContext.sql("SELECT * FROM Web_Session_Log limit 10")
Programmatically Specifying the Schema :
from pyspark.sql import *
sqlContext = SQLContext(sc)
lines = sc.textFile("/mnt/weblog.csv")
parts = lines.map(lambda l: l.split("\t"))
Web_Session_Log = parts.map(lambda p: (p[0], p[1],p[2], p[3],p[4]))
schemaString = "DATETIME USERID SESSIONID PRODUCTID REFERERURL"
fields = [StructField(field_name, StringType(), True) for field_name in
schemaString.split()]
schema = StructType(fields)
schemaPeople = sqlContext.createDataFrame(Web_Session_Log, schema)
schemaPeople.registerTempTable("Web_Session_Log")
results = sqlContext.sql("SELECT USERID FROM Web_Session_Log")
names = results.map(lambda p: "Name: " + p.USERID)
for name in names.collect():
print name
```

How to Run This Python Code:
First, go to pyspark prompt by running ./pyspark from your installation directory.
```
[root@ip-10-229-95-146 bin]# ./pyspark
```

You will see the following prompt:

```
Using Python version 2.6.6 (r266:84292, Jan 22 2014 09:42:36)
SparkContext available as sc, HiveContext available as sqlContext.
>>>
```

You can just launch pyspark from Spark installation bin directory and paste above code. Or, you can create a file with the above content as say mypyspark.py and run as

```
    ./bin/pyspark /yourdirectory/myspark.py
```

Analyze the result. Notice that the table used in this step was created in step 4. If you have created your own table, refer to that. If you run into any issue, you will need to debug. This will help you to learn what each line of code is doing.

You should see the following result

Step-7. Caching tables and Un-caching tables;

Caching tables
This is extremely useful when you are joining tiny dataset with huge dataset.

CACHE TABLE and UNCACHE TABLE statements are available to do the above making it very easy.

```
    To Cache a table :
    CACHE TABLE logs_last_month;
    To UnCache a table :
    UNCACHE TABLE logs_last_month;
```

Once as table is cached, you can use in your spark queries.

You should see the following result

Step-X. Accessing through Java database connectivity (JDBC)

Starting a JDBC server
Spark SQL provides JDBC connectivity, which is useful for connecting business intelligence (BI) tools to a spark cluster and other users or applications.

Here are the steps to launch a the Spark SQL JDBC server

From Spark Installation folder
```
cd /usr/lib/spark

./sbin/start-thriftserver.sh —master sparkMaster
```

You should see the following result

Step-X. Let's connect to the above server(JDBC) from Step-2 through Beenline:

Connecting to a Database

From Spark Installation folder
```
cd /usr/lib/spark

//10000 will be the default thrift server port

./bin/beenline -u jdbc:hive2://localhost:10000

0: jdbc:hive2://localhost:10000> show tables;
```

You should see the following result

What you should have learnt

1. What are the various components of Apache Spark?
2. What is the difference between Spark-shell and Spark-sql?
3. What is a RDD?
4. Why Spark is faster than Map-reduce FW?
5. How will you use Spark from your visualization tools?

Troubleshooting

If you get an exception looking like "ERROR SparkContext: Error initializing SparkContext. java.net.UnknownHostException:…" make sure you have the make of our computer added to the /etc/hosts file. For example add the line "127.0.0.1 <myhost>", where <myhost> is the name of our computer.